



Oracle Fusion AI Platform and Data Security Principles



Oracle Fusion AI

Jan, 2026 | Version [3.8]
Copyright © 2026, Oracle and/or its affiliates
Public

Purpose Statement

This document provides an overview of the Oracle Fusion AI Platform, its architecture, data flow and data security principles that are the foundation of Oracle Fusion AI for Fusion Cloud Applications, now updated to include generative AI capabilities. The information in this document is intended solely to help the reader understand basic functional, technical, and security aspects of Oracle Fusion AI to plan for business and IT objectives.

Disclaimer

The information contained herein is subject to change without notice and is not warranted to be error-free. If you find any errors, please report them to us in writing.

This document in any form, software or printed matter, contains proprietary information that is the exclusive property of Oracle. Your access to and use of this confidential material is subject to the terms and conditions of your Oracle software license and service agreement, which has been executed and with which you agree to comply. This document and information contained herein may not be disclosed, copied, reproduced, or distributed to anyone outside Oracle without prior written consent of Oracle. This document is not part of your license agreement, nor can it be incorporated into any contractual agreement with Oracle or its subsidiaries or affiliates.

This document is for informational purposes only and is intended solely to assist you in planning for the implementation and upgrade of the product features described. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described in this document remains at the sole discretion of Oracle. Due to the nature of the product architecture, it may not be possible to safely include all features described in this document without risking significant destabilization of the code.

Oracle and Java are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

This software and documentation may provide access to or information about content, products, and services from third parties. Oracle Corporation and its affiliates are not responsible for and expressly disclaim all warranties of any kind with respect to third-party content, products, and services unless otherwise set forth in an applicable agreement between you and Oracle. Oracle Corporation and its affiliates will not be responsible for any loss, costs, or damages incurred due to your access to or use of third-party content, products, or services, except as set forth in an applicable agreement between you and Oracle.

Introduction

Oracle Fusion AI is a family of embedded artificial intelligence (AI), machine learning (ML) and generative AI (GenAI) features within the Oracle Fusion Cloud Application suite that help customers get business insights faster, automate routine tasks, access personalized recommendations, and improve business performance.

Oracle Fusion AI features are pre-built by Oracle developers and do not require customer (user) data science or machine learning expertise to use. Oracle Fusion AI features within Oracle Fusion Cloud applications can be enabled simply and quickly by the customer, and do not require outside system integrator assistance to set up. Each AI feature can be easily enabled or disabled (see Oracle Help Center documentation for setup and enable/disable instructions for an [individual Fusion AI feature](#)).

The Oracle Fusion AI Platform, described within this document enables and orchestrates the use of Fusion AI features for Oracle Fusion customers.

To assist reader comprehension of this document, it is helpful to distinguish two major classes of AI that are processed by the Oracle Fusion AI Platform – classic AI and generative AI. Classic AI are essentially all non-generative AI features, and typically include use-case specific ML models. Examples of classic AI applications might be data classification, making predictions, and detecting anomalies, to name a few. Generative AI features utilize large language models (LLMs) and are adept at content creation, including creation of text, images, videos and even code. Common text creation applications include assisted authoring, document summarization, and suggestions.

Details of LLMs used within Fusion SaaS Applications can be found on [MOS KB796229](#).

Classic AI features use machine learning algorithms trained on individual customer data and learn over time through customer use to provide increased value and immediate customer benefits. GenAI features use pre-trained LLMs and are not trained on customer data.

Each Oracle Fusion Cloud Application customer is provisioned with a securely partitioned Oracle Fusion AI Platform environment, within their local OCI region. The Oracle Fusion AI Platform is designed for high performance and scalability, while conforming to Oracle's stringent data security principles.

The Oracle Fusion AI Platform is a multi-tenant environment, where each customer's data and trained ML models (for classic AI) remain completely isolated. No customer data is shared, mixed, or combined to train ML models, or to process GenAI requests.

Pre-trained generative AI models (LLMs) are securely accessed by the Oracle Fusion AI Platform within a secure Oracle Cloud Infrastructure (OCI) environment (which may be outside of the immediate local OCI region). However, please note that in the case of special OCI realms such as EUSC (European Union Sovereign Cloud) then AI traffic will route only to the appropriate OCI region, e.g. Frankfurt. Data is not persisted in the LLMs, i.e., after a prompt is processed, no customer data persists, and prompt data does not train the LLM. All data processed within the LLM is completely isolated, and no customer data is shared or combined.

AI outcomes, predictions, and Gen AI responses are sent from a customer's Fusion AI Platform environment to the customer's Oracle Fusion Application via secure REST APIs via typical business workflows.

This document will describe the architecture, data flow, and security principles of the Oracle Fusion AI Platform, for classic AI and GenAI, working with Oracle Fusion Applications.

Structure of Fusion Applications with the Oracle Fusion AI Platform within OCI

Oracle Fusion AI, leveraging the Oracle Fusion AI Platform, is a composite service of Oracle Fusion Cloud Applications designed to run multiple AI use cases that serve the Oracle Fusion Cloud Application Suite. Oracle SaaS developers and data scientists design and build machine learning pipelines (AI application features) which produce models that are trained with data ingested from Fusion applications and deliver predictions back to Fusion databases or user interfaces.

For Gen AI features, Oracle SaaS developers and data scientists create prompts, through prompt engineering, that are securely sent to the Oracle Generative AI Service for processing with results returned to the Oracle Fusion Application. The SaaS developers then manage the lifecycle of the AI application features.

The basic structures utilized to deploy Oracle Fusion AI are summarized as follows:

- Oracle Fusion AI Platform for Oracle SaaS developers built on Oracle Cloud Infrastructure (OCI) and Oracle Cloud Infrastructure Data Science (OCI-DS).
- Pipelines leveraging trained ML models to process and respond to requests for specific user workflows (classic AI).
- Pipelines delivering pre-engineered prompts developed by Oracle SaaS developers and Oracle data scientists to process and respond to requests for specific user workflows (Gen AI).
- OCI Data Stores for storage of isolated customer data used for ML model training, customer-specific trained ML models, and metric data used for feature evaluation, performance, and support.
- Oracle Generative AI Service including Oracle Dedicated AI Cluster (DAC) and LLM.




Figure 1. Depiction of basic structures utilized to deploy Oracle Fusion AI features for classic AI and Generative AI (Including Agentic AI).

Functional area	Description
BOSS Services	Business Object Spectra Services provides a set of REST APIs for accessing and manipulating data within Fusion Applications. Used by

	Fusion AI to write back to Fusion DB, an example being writing metrics back to Fusion DB tables in customer tenancy
ML Prediction Service / ML Apps (Classic AI)	Handles workflow for Classic AI aka Predictive AI. Concerned with data ingest and preparation, model training and deployment, and prediction serving.
Gen AI Orchestration	Responsible for consuming a request containing the context of the use case and customer data fields from Fusion and making the call to the relevant AI services (Gen AI, Q&A, Agentic, routing, guardrails, etc.) and responding back to Fusion.
Fairness Service	Determines if request should be allowed or denied, for example if one AI feature is putting excessive load at the expense of other AI features being used.
RAG Q&A	An Agent that manages RAG (Retrieval Augmented Generation) flow including document ingestion, chunking, generation of embeddings as well as managing question and response generation.
Agent Framework	Workflow engine that manages Agent flows, including reasoning, tool execution (RAG, FA REST Object, Calculator, Email, Deeplink, etc.)
Metrics Scoring / Publishing	Evaluates responses from LLM, scores and records them into customer store for future reporting capability for customer use.
Routing Service	Determines how to route LLM calls taking into account data jurisdiction rules. Please review Fusion Generative AI Data Center Regional Availability (MOS KB796299) for further information.
OCI Guardrails	Guardrail service which has scanners for Prompt Injection, Content Moderation, etc. The Guardrail service is utilized by the embedded Generative AI features within Fusion SaaS Applications with each AI feature determining which specific guardrails are required.
Oracle Search	Current repository for vector storage as required in RAG (Retrieval Augmentation Generation) aka Q&A Use Case. Will be replaced by 23ai Vector Search feature on customer's Fusion SaaS DB in due course.
OCI Region (Chicago, Frankfurt)	The Large Language Models (LLM) utilized by Fusion Applications are currently located in Chicago and Frankfurt. Individual regions, e.g. Ashburn, London, Japan, etc. are mapped to one of these two regions for serving Generative/Agentic calls. Please see detail on Routing Service above. Important: For OCI realms such as EUSC (European Union Sovereign Cloud) then AI traffic will route only to the appropriate OCI region, e.g. Frankfurt.
Cohere Cmd-R	Command-R LLM from Cohere which manages Generative AI features of Author, Summarize, Extract, Compare for the embedded AI features (not Agent Studio).
Llama	Meta LLM which is used for RAG (Retrieval Augmented Generated aka Q&A) use case.
Cohere Embed	Used to generate embedding representations for insertion into a vector store and is used in RAG use cases.
OpenAI GPT-4.1 mini / GPT-5 mini	Used in Agentic Flows. Please note that OpenAI is hosted outside of OCI Tenancy and is accessed via a passthrough service in the OCI tenancy. OpenAI use is subject to ZDR (Zero Data Retention) policy,

	meaning no data is retained on the OpenAI tenancy. Llama can be used in Agentic AI flows in place of OpenAI GPT if required. Please note that GPT-4.1 is currently not available within EUSC and certain other OCI realms outside of OC1 (commercial realm).
GPT-OSS	Used in Agentic Flows. Please note that this model is hosted with OCI.
Object Store Customer Bucket	Tenant-specific OCI storage bucket, used for Classic/Predictive AI as well as storage of execution and evaluation metrics for Generative AI

Logical Architecture of Oracle Fusion AI Platform Interacting with Fusion Applications (Classic AI)

The Oracle Fusion AI Platform for classic AI, can be described by logical components and data flows.




Figure 2. Logical Architecture of the Fusion AI Platform (Classic) and integration with Fusion Applications (Oracle Cloud Applications).

Figure 2 illustrates the logical architecture and data flow within the Oracle Fusion AI Platform and as it communicates to an Oracle Fusion Cloud Application for classic AI use cases.

1. Data Ingest from the Oracle Fusion Cloud Application

- Oracle Cloud and Fusion AI development teams specify data to ingest for a specific feature.
- Adhering to the principle of data minimization, ingested data is limited to what is specifically required for ML model training for a particular AI feature and is pre-approved for ingestion through Oracle CSSAP (Corporate Security Solution Assurance Process) security review procedures.
 - CSSAP is a security review process developed by Corporate Security Architecture, Global Information Security, Global Product Security, Oracle Global IT, and Oracle's IT organizations to provide comprehensive information-security management review. (For more information, please see <https://www.oracle.com/corporate/security-practices/corporate/governance/security-architecture.html>).
- No human readable personal identifiable information (PII) is ingested.
- An initial bulk data ingestion is performed at setup time, then data is ingested incrementally at a periodicity appropriate for a particular use case.
- Data is written to OCI Object Store within a tenant-specific bucket (i.e., per customer).

2. Data Preparation follows to prepare the data for model training.

- Multiple objects that have been ingested are read from storage and prepared for training purposes.
- The prepared data is written back to OCI Object Store.

3. Model Training and Deployment

- Data is read from OCI Object Store, and a training flow is triggered.
- The trained model is deployed, and prediction endpoints are made available.
- Data flows are performed for model retraining, with monitoring for data and model quality.

- The services within the OCI Data Science Service have isolated infrastructure.

4. Prediction Serving

- The Fusion Application calls a Fusion AI endpoint that routes to the appropriate prediction REST API for the customer.

Logical Architecture of Oracle Fusion AI Platform Interacting with Fusion Applications (Generative AI)




Figure 3. Logical Architecture of the Fusion AI Platform (Gen AI) and integration with Fusion Applications (Oracle Cloud Applications).

1. Generative AI Request from the Oracle Fusion Cloud Application initiates the request to the LLM.

- Once the Generative AI feature is enabled, invoking Generative AI happens in one of two ways:
 - a) The user initiates by providing user input and requesting assistance from AI through the application UI.
 - b) The application has enough context to issue an invocation implicitly, based on the activities of the user.

The invocation step sends the pre-engineered prompt along with the relevant customer-specific context to the Fusion AI Platform Orchestration Service through a REST API service.

- Prompt design, engineered prompts, and data features selected for a particular generative AI use case are developed to mitigate response “hallucinations” and provide relevant, useful responses. Prompts and prompt features are pre-approved for use through Oracle CSSAP security review procedures.
 - CSSAP is a security review process developed by Corporate Security Architecture, Global Information Security, Global Product Security, Oracle Global IT, and Oracle's IT organizations to provide comprehensive information-security management review. (For more information, please

see <https://www.oracle.com/corporate/security-practices/corporate/governance/security-architecture.html>).

- Prompts are engineered not to include human readable personal identifiable information (PII), as specified by the CSSAP guidelines. However, use cases that allow for user input may include PII data as part of the customer input.
- 2. **Fairness Service** determines if request should be allowed or denied, for example if one AI feature is putting excessive load at the expense of other AI features being used. It looks at current load on the LLM cluster along with recent usage for a given agent / tenant to determine if the request is allowed or if it needs to be throttled back. Goal is to serve as many requests as possible without compromising stability of the service. Fairness service is an internal framework control which is not modifiable by the customer.
- 3. **Routing Service** determines how to route LLM calls taking into account data jurisdiction rules.
- 4. **Prompt Assembly** constructs prompt using prompt template, customer data, guard rail requirements, LLM parameters.
- 5. **Prompt Response** sends the response from the OCI Generative AI Service, i.e., Dedicated AI Cluster (DAC) to the Guardrail Service.
 - Neither the engineered prompt, nor the customer data included for context are persisted in the OCI Generative AI Service. No customer data is used to train a shared generative AI model.
- 6. **Guardrail Service** evaluates the LLM response for compliance to use case-specific requirements. Metrics that the Guardrail Service evaluates include, but are not limited to, accuracy, toxicity, sentiment, and repetitiveness. The response from the Guardrail Service is pass/fail.
- 7. **Metric Store** is used for monitoring and support of the OCI Generative AI Service. The prompts, responses, guardrail responses, and usage metrics are initially stored in an Autonomous Transaction Processing (ATP) Gen AI Metrics Store for a short duration (maximum 28 hours). The metrics data is moved for long term storage to the Metrics Store in OCI Object Store which is partitioned by tenancy.
- 8. **Response Delivery** sends the completed LLM response to the Fusion application via REST API. Additionally, the response can be a “fail” from the Guardrail Service, whereby the Fusion application may attempt a retry.

Other notes related to the Fusion AI Platform architecture and workflow for both classic AI and generative AI:

- OCI Object Store is the final destination for usage metrics.
- IDCS (Oracle Identity Cloud Service) and WTSS (Web Tier Security Service) are used for protecting REST Endpoints and Authorization.
- The OCI Generative AI Service may be located in an OCI region outside of the customer’s local OCI region. Data transport between OCI regions is encrypted and secured through TLS 1.2. or TLS 1.3 depending on customer configuration. Please review [MOS KB Article #173925](#). Please note that that for OCI realms such as EUSC (European Union Sovereign Cloud) then AI traffic will route only to the appropriate OCI region, e.g. Frankfurt.
- No data is stored or persisted in the OCI Generative AI Service, and no customer data is used to train LLMs.

Guardrails

Guardrails for Fusion SaaS Applications are one or more of detection based, Instruction based, LLM based.

Detection based guardrails are used to detect and prevent prompt injection attacks as well as guard against toxic and other inappropriate content for both the prompt and response (content Moderation). These detection based guardrails are in place for the embedded Generative AI features within Fusion SaaS Applications.

Content Moderation is part of the OCI Guardrails Service and is divided into the following four categories:-

- Hate and harassment, such as identity attacks, insults, threats of violence, and sexual aggression
- Self-inflicted harm, such as self-harm and eating-disorder promotion
- Ideological harm, such as extremism, terrorism, organized crime, and misinformation
- Exploitation, such as scams and sexual abuse

Instruction based guardrails are instructions which form part of the system prompt which help guide the LLM from giving undesirable responses by responding to off-topic requests. Typically this form of protection is used (and implemented) by Agents which are deployed from Fusion AI Agent Studio.

LLM based guardrails are safeguards implemented by the third-party LLM providers to protect the Model at both the prompt and the response level. Such safeguards can act on inappropriate prompting as well as ensuring that the response from the model is safe. Please review model cards and responsible AI statements from the various providers for more detail, e.g. [Llama 3 AI Responsibility](#), [OpenAI GPT-5 System Card](#).

RAG (Retrieval Augmented Generation) Agent / Answer a Question

A subset of Author/Summarize use case is the ability to answer user questions on a given topic and for Gen AI to base its answers on actual corporate information related to the question. For example, in the use case of Benefit Advisor an employee may pose a question, “What health insurance plans are available to me and my family”. In order for GenAI to answer that question accurately, it needs access to the relevant benefits policy documents for the organization that the employee works for. The RAG Agent, deployed as part of our Agentic AI architecture, enables this type of use case.

In the above example, Benefit Policy documents would be loaded into Fusion Applications by the Benefits Administrator and processed by the system via a concurrent process. The overall flow for document ingestion is as follows:

1. Write a copy of the document(s) into the tenants document store (UCM – Universal Content Management)
2. The file processing service pulls the document(s) from UCM converts it into to suitable format ready for chunking process.
3. The chunking process parses each document and splits it into chunks using an algorithm to determine optimal chunk size / split according to document headings, sections, etc.
4. Processed chunks and information about the source document are stored within a set of FA tables.

5. Embeddings (vector representations) are computed for each chunk via calls to the Embeddings LLM (e.g. Cohere Embed Multilingual)
6. Embeddings are stored within the vector store and indexed into a knowledge base which will be used for answering user questions.

Once processed, the policy documents are available to the Benefits Policy Advisor to assist it in answering employee questions during a Q&A session.

When the employee poses a question, the following occurs:

1. User question is converted into a formatted query suitable for processing.
2. The formatted query is converted to a vector (via same embedding engine used to process the underlying policy document).
3. The knowledge base is queried (vector similarity search) to retrieve semantically and lexically similar chunks of the relevant policy document.
4. The retrieved chunks are ranked according to the nearest match to the user query.
5. An LLM Prompt is assembled consisting of the question and relevant policy document chunks. The document chunks provides context to the LLM when considering its response. This combined Prompt helps the LLM to produce a more reliable and relevant response by ensuring the answer is based on factual information per the user question originally asked.
6. The fully formed Prompt (question and context) is sent to the LLM to generate a response to the user query.
7. Links (citations) to the relevant parts of the policy document are created.
8. The user is presented with the response along with the relevant citations. The citations reassures the user that the generated answers come from a reliable source, (i.e. relevant policy document) and provides ability to review the relevant section of the policy document for further information.

It is **important** to note that the LLM is not trained on the uploaded policy documents. These source policy documents are stored within the customers own document store in their Fusion instance and the embeddings (representations of the content of the policy document to be used in semantic search) reside within the customers stripe in the Spectra tenancy. No data is shared between customers. Figure 4 outlines the basic architecture of this approach.




Figure 4. Logical Architecture of the RAG Use Case (will change in future releases)

ORACLE AI AGENT STUDIO

A broad outline of the Fusion Applications Agentic AI Architecture can be found in the diagram below.




Figure 5. Depiction of basic architecture of Fusion Applications Agent Studio and Fusion AI Agentic AI.

Detail of major sections in the diagram above is as follows:

Fusion AI Agent Studio: Low Code development platform to design and configure Agentic Workflows (Agent Configuration). Enables users (admins) to:-

- Instantiate a new Agent Team from an Oracle shipped templates with minor or no modification.
- Instantiate a new Agent Team by extensively modifying a seeded Agent Template thereby creating a new “Custom” Agent Team
- Create brand new Agent Team (“Custom” Agent Team)
- Test and Evaluate Agent Teams prior to deploying for production use
- Version, Save and Deploy Agent Configuration to customer’s FA DB

Fusion AI Agentic Services: Spectra/Hybrid

- Spectra: Customer’s OKE (Oracle Kubernetes Engine) Container that has Fusion AI Services deployed
- FAI Adjunct Tenancy: Regional OCI tenancy which runs certain Fusion AI services as well as being the repository for the Object Store (customer bucket) which is used in both Classic and some GenAI use cases.
- OCI Resources: DAC cluster containing LLMs (e.g. Cohere, Llama, etc) and FA Agentic AI resources as well as pass through service for using externally hosted OpenAI LLM.
- LLMs perform varied tasks such as:
 - Convert initial call to Agent Workflow to a plan which can be executed by the Execution Framework
 - Convert document chunks to vector for RAG use case
 - Author, summarize, text for Generative AI use case
- Read Agent Definition: When Agent Team is called from the Fusion Applications Layer, the Agent Configuration is initially read from the FA DB and converted to framework elements required to be executed by the Execution Framework

Fusion Data Layer: The customer’s Fusion SaaS data layer consisting of :-

- Fusion Applications Database: Transactional data (e.g. HCM, ERP, SCM, etc.), meta data (e.g. Agent Configuration), logs (e.g. AI Metrics Store)
- UCM (Universal Content Manager): Document store which is used in RAG (Retrieval Augmented Generation) use case.
- Object Store (Customer Bucket): Customer partitioned data store containing mix of data including Classic AI, Metrics, etc., held within the Fusion AI Adjunct Tenancy

Fusion Application Layer: Deploying an AI Agent Team means that it is available to be invoked as an endpoint which can be called from the Fusion Application Layer, typically this would be via:

- UI (e.g. user starts a chat session and asks a Benefits Enrollment question)
- Fusion backed process (initiated via an API call to the Agent passing the Agent Team Code).

Customer Partitioning

The Oracle Fusion AI Platform can support many customers for the same AI feature(s) – a fundamental requirement for cloud applications.




Figure 4. Fusion AI Platform multi-tenancy architecture showing customer isolation for classic and generative AI.

- Oracle makes use of core OCI technologies to provide tenant isolation and separate ML models per customer, i.e., training and prediction flows are per customer. No customer data is shared or pooled for ML training, and no customer data is used to train LLMs.
- Customers are isolated throughout the pipeline using techniques approved by Oracle's CSSAP security review process.
- Isolation is achieved through use of Kubernetes namespaces (OKE is the Oracle Container Engine for Kubernetes), customer tenant-specific OCI Virtual Machines (VMs) and corresponding separate storage buckets for each tenant, which are only accessible by the specific namespace or VM.

Data Science Monitoring

For Oracle to support a robust AI solution for multiple customers, data science monitoring is required. Data science monitoring ensures proper functioning for the following:

- Model training, tuning, and retraining (for classic AI)
- Generating predictions (for classic AI)
- Generating responses from prompts (for GenAI)
- Logging metrics and classic AI data model (prediction) quality, and GenAI response quality
- Automating responses to quality issues

To support the above functions, varying types of metrics data are required to be gathered and examined by Oracle data scientists, including ML model metrics, prediction metrics, and generative AI response metrics. In addition, monitoring system telemetry (performance and latency metrics), data quality, model quality, and prediction quality are fundamental for supporting robust AI solutions for multiple customers.




Figure 5. OCI Data Science Workbench for monitoring of AI functions and performance.

- Oracle Data scientists have access to the metrics stored within the OCI Object Store through the Data Science Workbench. It is important to note that metrics available for ongoing data science monitoring do not include customer data.
- Data at rest (e.g., in OCI Object Store) is encrypted with AES-256. In-transit control plane data is encrypted by using Transport Layer Security (TLS) 1.2 or later.
- In order to deliver a robust SaaS offering, certain circumstances may require that Oracle Data Scientists obtain restricted access to customer data. Access follows the approved and secure OCI policy rules in place as for other operational data access. In the case of EUSC (European Union Sovereign Cloud) then such access is restricted to authorized personal within region.

AI Governance & Standards Accreditation

Oracle participates in a number of [standard setting bodies](#) and related organizations, detail of which may be helpful to the reader. The [Cloud Compliance](#) portal lists out attestations that Oracle has achieved.

Cloud Security Alliance (CSA)

Oracle completed the AI CIAQ (Consensus Assessments Initiative Questionnaire) which has been made available in the STAR (Security, Trust, Assurance, and Risk) Registry. This resource will be helpful to the reader of this white paper since it contains questions on AI Governance and related topics.

- [Oracle AI CIAQ Self-assessment](#)
- [STAR Registry Listing for Oracle SaaS Services](#)
- [STAR Certification \(CCMv4\)](#) - Level 2 Certification Cloud Controls Matrix

International Organization for Standardization (ISO)

There are a number of ISO standards which Oracle is engaged with, see [Cloud Compliance](#) portal for list. An important new standard in the Artificial Intelligence space is [ISO 42001:2023](#) which specifies requirements for establishing, implementing, maintaining, and continually improving an Artificial Intelligence Management System (AIMS) within organizations.

ISO/IEC 42001:2023 - Information Technology - Artificial Intelligence - Management System

- Issued [14681141-1](#) certification of the standard for AIMS (AI Management System) supporting SaaS Applications
- Issued [1250166-1](#) certification of the standard for AIMS (AI Management System) supporting OCI Infrastructure AI Services

Additional Information on OCI Realms and applicability to Fusion SaaS AI Architecture

This white paper has been concerned with describing the Fusion AI Architecture as it pertains to OCI Commercial Realm, otherwise known internally as OC1. There may be some differences on the other realms usually concerning model availability and location, you can find further information on the following MOS KB articles:-

- [KB7996229 – What Large Language Models \(LLM\) are in use or supported by Fusion SaaS Applications](#)
- [KB796299 – Fusion Generative AI Data Center Region Availability](#)

CONNECT WITH US

Call +1.800.ORACLE1 or visit [oracle.com](#).

Outside North America, find your local office at [oracle.com/contact](#).

[blogs.oracle.com](#)

Copyright © 2026, Oracle and/or its affiliates. All rights reserved. This document is provided for information purposes only, and the contents hereof are subject to change without notice. This document is not warranted to be error-free, nor subject to any other warranties or conditions, whether expressed orally or implied in law, including implied warranties and conditions of merchantability or fitness for a particular purpose. We specifically disclaim any liability with respect to this document, and no contractual obligations are formed either directly or indirectly by this document. This document may not be reproduced or transmitted in any form or by any means, electronic or mechanical, for any purpose, without our prior written permission.

Oracle and Java are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

Intel and Intel Xeon are trademarks or registered trademarks of Intel Corporation. All SPARC trademarks are used under license and are trademarks or registered trademarks of SPARC International, Inc. AMD, Opteron, the AMD logo, and the AMD Opteron logo are trademarks or registered trademarks of Advanced Micro Devices. UNIX is a registered trademark of The Open Group. 0120

Disclaimer: This document is for informational purposes. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, timing, and pricing of any features or functionality described in this document may change and remains at the sole discretion of Oracle Corporation.